

nlint: nginx configuration file linter written in Haskell

Federico Bond

February 15, 2016

Contents

1	Introduction	1
2	Syntax definition	1
3	Design Considerations	2
4	Difficulties encountered	2
5	Preliminary research	3
6	Future extensions	3
7	Futher research	3
8	Conclusions	3
	Bibliography	4

1 Introduction

The following report describes the design of a linter for **nginx** configuration files, written in Haskell. **nginx** is a widely used HTTP and reverse proxy server. The motivation for this project was to further the author's understanding of functional parsing techniques, and apply them in a real-world context.

2 Syntax definition

The **nginx** configuration uses a C-like syntax consisting of directives, which let you specify the behaviour of the server, and blocks, which define contexts in

which these directives apply.

A directive consists of an identifier followed by a list of arguments and a semi-colon.

A block consists of an identifier followed by an optional list of arguments and a list of blocks or directives surrounded by braces.

An example configuration file is shown below:

```
user www-data;
worker_processes auto;
pid /var/run/nginx.pid;

events {
    worker_connections 2048;
    multi_accept on;
    use epoll;
}

http {
    # default http server
    server {
        listen 80;
        root /srv/website;
    }
}
```

3 Design Considerations

The project has two main components, a parser and a linter. The parser is built with the `parsec` (Daan Leijen 2001) library of monadic parser combinators. The linter uses reverse function composition to apply different filters to the AST generated by the parser, finally returning a list of problematic declarations which is then mapped to individual lint messages.

An additional library, `ansi-terminal`, was used for colored console output.

4 Difficulties encountered

The most difficult part of this project was to find a proper abstraction for constructing linter rules, which avoided unnecessary boilerplate. The author considered the use of Monads, Zippers and Traversable, but finally settled for simple composition of filter functions, which proved good enough for the task at hand and kept the mental overhead at a reasonable level.

The error output from `ghc` also proved hard to understand at first, for someone accustomed to traditional stack traces or syntax errors in imperative / object-oriented code, but with practice they became quite natural.

5 Preliminary research

The research for this project included reading the original `parsec` paper, several articles and tutorials on Monads and the documentation of more than a dozen other libraries.

Other topics studied, whose understanding did not contribute much to this project, were Functional Reactive Programming, Arrows and the Template Haskell extension to GHC.

6 Future extensions

Throughout the duration of the project, the author kept a To Do list with ideas for future extensions. A sample of them has been reproduced below:

- Add rules which depend on the presence or absense of siblings or otherwise related declarations.
- Add multiple output formats for the linter messages, to ease integration with text editors and other tools.
- Add command line options for ignoring certain rules.
- Ignore linter errors when preceded by a specially formatted comment.
- Use the QuickCheck library for testing the codebase

7 Futher research

The author would like to explore more idiomatic alternatives for constructing linter rules. The `hlint` and `shellcheck` projects could prove useful for this task.

The author would also like to deepen his understanding of Monad Transformers, which appear useful for cleaning up the highly nested code in the `main` function.

8 Conclusions

The author of this project had never written a complete program in Haskell before. The experience proved highly satisfactory and encouraged him to use Haskell and other functional programming languages like Elm and PureScript in future scenarios.

Having said that, the state of documentation in the Haskell ecosystem leaved much to be desired. The author found the need to read academic papers to understand how to use each library a bit tedious, but grew accustomed to it over time. Many libraries could benefit from a concise README with sample code and common caveats, specially considering that the heavy use of types in Haskell makes it easy to find what you need from there.

Bibliography

Daan Leijen, Erik Meijer. 2001. “Parsec: Direct Style Monadic Parser Combinators for the Real World.” <http://research.microsoft.com/en-us/um/people/daan/download/papers/parsec-paper.pdf>.